

Image Manipulation Using Matrix Techniques

Names	Student ID	Professor	TA	Recitation
Will Farmer	101446930	Mimi Dai	Amrik Sen	608
Jeffrey Milhorn	100556107	Kevin Manley	Ed Yasutake	605
Patrick Harrington	100411000	Mimi Dai	Ash Same	618

Friday, March 22

Contents

Table of Contents	1
1 Reading Image Files & Grayscale Conversion	2
2 Horizontal Shifting	4
3 Vertical Shifting	4
4 Inversion	7
5 Transposition	7
6 DST	7
7 Restrictions on Compression with the Discrete Sine Transform	8
8 Compression	9
9 Optimization	9
10 Conclusion	10
11 Code	10
11.1 Python	10
11.2 MATLAB Code	17
11.3 MATLAB Code	18

List of Figures

List of Figures	1
1 Provided Images	2
2 A simple RGB image	2
3 A given image split into its three primary color channels	3
4 A given image split into its three primary color channels, but only intensity of each color is shown. . .	3
5 Grayscale Images	4
6 Horizontally Shifted Images	5
7 Vertically Shifted Images	5
8 An example of a transposed image	7
9 Δ^2 of $S(n)$	8
10 The plot of the Discrete Sine Transform	8

Introduction

Since images stored on computers are simply matrices where each element represents a pixel, matrix methods learned in class can be used to modify images. The purpose of this project was to apply matrix manipulations on given image files, shown below as Figure 1a and Figure 1b.



(a) Photo 1



(b) Photo 2

Figure 1: Provided Images

1 Reading Image Files & Grayscale Conversion

Colored images have an interesting, although problematic property; they do not readily lend themselves to matrix manipulation because in order to get color images, separate values are used to represent each primary color, which are then mixed together for the final color. For example, in Figure 2, the block represents very simple a 2×2 pixel image.




Figure 2: A simple RGB image

This very simple image can be represented as either a trio of primary color matrices where each entry in each primary color matrix corresponds to the same pixel:

$$\underbrace{\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}}_{\text{Red Matrix}}, \underbrace{\begin{bmatrix} 0 & 0 \\ 1 & 1 \end{bmatrix}}_{\text{Blue Matrix}}, \underbrace{\begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix}}_{\text{Green Matrix}}$$

A single matrix may be used, with each entry being a submatrix, wherein each element in the submatrix corresponds to a primary color.

$$\begin{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} & \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} \\ \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix} & \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} \end{bmatrix}$$

Using one of the given images, the splitting of color channels gives the following set of images shown in Figure 3.



Figure 3: A given image split into its three primary color channels

While it is possible to manipulate color images, it would be far simpler to manipulate *grayscale* images, where only the final intensity is concerned. To do this, each color is considered independently for its intensity alone as shown in Figure 4, where it may be scaled, and then added together to produce a final black-and-white image, which is a matrix where each entry is a single value. Note how the third panel representing the blue color channel is darker – this implies that blue is a less intense color in the image.



Figure 4: A given image split into its three primary color channels, but only intensity of each color is shown.

Since each primary color is freely editable, it is simple to scale the intensity of each before mixing; in our report, we used 30% of the red channel, 59% of the green channel and 11% of the blue channel. The final outputs for both given images can be seen in Figure 5. Note how the final output is lighter than any of the individual color channels.



(a) Photo 1 - Grayscale



(b) Photo 2 - Grayscale

Figure 5: Grayscale Images

2 Horizontal Shifting

Now that we are working in grayscale, it is far more straightforward to manipulate aspects of the image, such as its horizontal position. Since we are dealing with a normal matrix, transforming the positions of columns requires only that we multiply the image matrix by a transformation identity matrix.

As discussed in the lab instructions, to shift an image horizontally without losing information requires the use of a transformation matrix as shown below.


$$\underbrace{\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}}_{\text{Identity Matrix}} \Rightarrow \underbrace{\begin{bmatrix} 0 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 0 & 0 \end{bmatrix}}_{\text{Transformation Matrix}}$$

$$\begin{bmatrix} 0 & 0 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix} \cdot \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} = \underbrace{\begin{bmatrix} c & a & b \\ f & d & e \\ i & g & h \end{bmatrix}}_{\text{The horizontally shifted matrix}}$$

3 Vertical Shifting

Very similar to the horizontal position change, the vertical position change merely requires the transformation matrix to be shifted row-wise as opposed to column-wise.

$$\underbrace{\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}}_{\text{Identity Matrix}} \Rightarrow \underbrace{\begin{bmatrix} 0 & 0 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix}}_{\text{Transformation Matrix}}$$



(a) Photo 1 Horizontal Shift




(b) Photo 2 - Horizontal Shift


Figure 6: Horizontally Shifted Images

Unlike the horizontal matrix shift, the order by which the transformation matrix is applied is reversed:

$$\begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \cdot \begin{bmatrix} 0 & 0 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix} = \underbrace{\begin{bmatrix} g & h & i \\ a & b & c \\ d & e & f \end{bmatrix}}_{\text{The vertically shifted matrix}}$$



(a) Photo 1 - Vertical and Horizontal Shift



(b) Photo 2 - Vertical and Horizontal Shift

Figure 7: Vertically Shifted Images

We can do both horizontal and vertical translations on our image matrix, but we must do the operations separately for photo1.jpg since they do not involve the same number of iterations. For example, we can first do the horizontal translation by using the same procedure above where the transformation matrix is second in the matrix multiplication (the transformation matrix would be dimension $n \times n$, where n equals the column dimension of photo1.jpg, 408). After

performing the 240 iterations of this horizontal translation we can then translate the image matrix vertically. We now place the transformation matrix first in the matrix multiplication; its dimensions must match the row dimension of the image. Therefore, this vertical transformation matrix is 201×201 .

4 Inversion

In order to flip a matrix upside down, we first had to generate an identity matrix of the appropriate size where the rows had the opposite diagonal direction.

$$\underbrace{\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}}_{\text{Identity Matrix}} \Rightarrow \underbrace{\begin{bmatrix} 0 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \end{bmatrix}}_{\text{Transformation Matrix}}$$

This was done by setting up the identity matrix as a two dimensional array; in other words, a list of lists. Then this list was iterated through and each list in the main list was flipped front-to-back. This action had the same effect as flipping the entire matrix on the horizontal axis. Finally, as before with the shifting process, we multiplied the matrix on the appropriate side of the matrix.

$$\begin{bmatrix} 0 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \end{bmatrix} \cdot \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} = \underbrace{\begin{bmatrix} g & h & i \\ d & e & f \\ a & b & c \end{bmatrix}}_{\text{The inverted matrix}}$$

5 Transposition

It is simple to visualize the effect of transposing a matrix; it would be a rotation about the main diagonal. The resulting image will be rotated 90° . Taking the transpose again would give the original image orientation following the properties of transposed matrices:

$$A = (A^T)^T$$

The effect can be seen in Figure 8:



Figure 8: An example of a transposed image

6 DST

From the plot of the determinant squared of S as a function of n for n from 1 to 32 shown in Figure 9, it can be seen that the determinant has strictly discrete values of either 1 or -1, and follows a sinusoidal pattern. It is also noticeable that the plot is an odd function.

The Discrete Sine Transform has the following equation:

$$S_{i,j} = \sqrt{\frac{2}{n}} \sin\left(\frac{\pi(i - \frac{1}{2})(j - \frac{1}{2})}{n}\right) \quad (1)$$

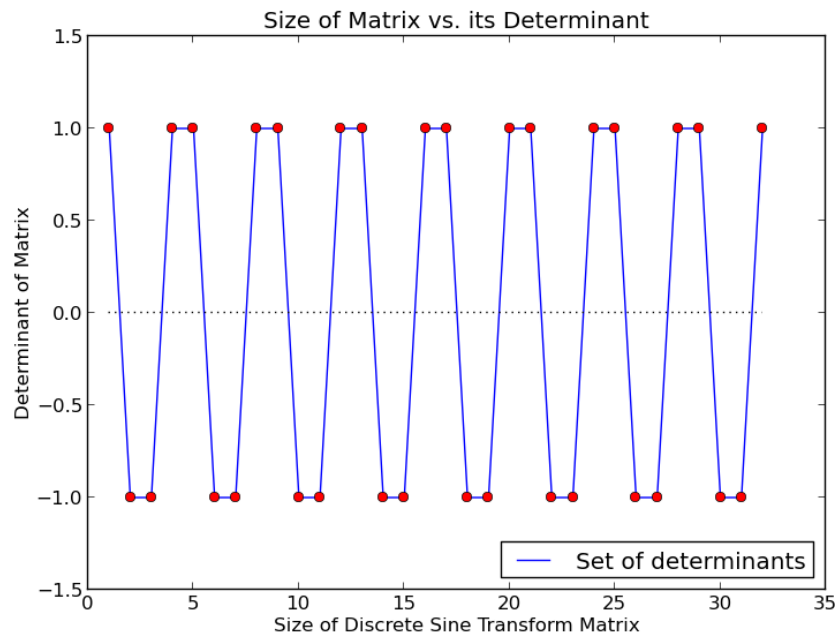
Figure 9: Δ^2 of $S(n)$ 

Figure 10: The plot of the Discrete Sine Transform

7 Restrictions on Compression with the Discrete Sine Transform

With the given equation to transform images using the Discrete Sine Transform (1), there does exist a limitation on the initial image aspect ratio – the image *must* be square. If it is not square, then the dot product will not work, and the image will not be compressed. The reason behind this is that since we are performing a dot product on the same matrix on either side, we know that in order for it to work it needs to be the same size after either operation is performed. The only matrix this holds true for is a square matrix.

That being said, the code below expresses a different algorithm. Instead of being limited to square matrices through the nuances of dot products, the code instead separates the two operations and performs them separately using two differently sized DST matrices. This algorithm is not limited by square matrices since it creates a new DST matrix for each operation.

```

1 def dst(image):
2     '''
3     If given a grayscale image array, use the DST formula
4     and return the result
5     Uses this method:
6         image = X
7         DST    = S
8         Y = S.(X.S)
9     '''
10    rows    = numpy.dot(image, create_S(len(image[0])))
11    columns = numpy.dot(create_S(len(image)), rows)
12    return columns

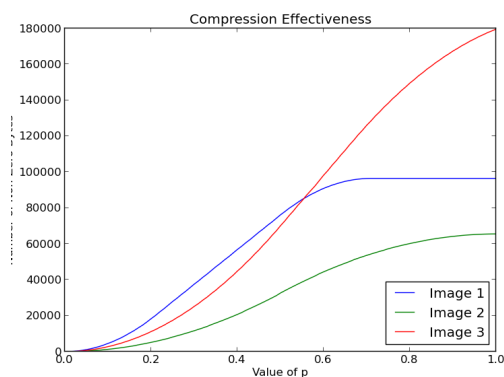
```

8 Compression


9 Optimization

Human vision has noticeable thresholds for the perception of light frequency. We are a lot more sensitive to lower frequencies compared to higher ones. JPEG image compression involves identifying high frequency pixel groupings and removing them from the image; less data in the image matrix means that it takes up less digital storage. We can compress the image by using the DST to identify high frequency data. We can also vary the extent of compression using a variable p , which goes from 0 to 1 where 0 represents a blank image and 1 represents an uncompressed image.

Because our equation focusses only on the high frequency values of the image, we can eliminate many pixels before our brains register a degradation in quality. Below is a graph that shows how much data is removed from each image as p gets larger.



(a) Bit Count vs. p Value



(b) Bit Ratio vs. p Value


Images can be compressed to low values of p without being noticeable, especially if the images are primarily uniform and consist of low frequency pixels. This is because of the threshold frequencies in human vision.

Due to the inherent nature of images having different amounts of high frequency values, different values of p will be appropriate for different images. For our first image, it initially will not have a large difference in quality as p gets smaller, however for our last image it will immediately start reducing in size. Therefore different values of P are appropriate for these different images.

10 Conclusion

Since digital images are represented as three dimensional matrices, image manipulation involves matrix operations. By using matrix multiplication, transpose, and the Discrete Sine Transformation (DST), we were able to translate, rotate, and compress multiple sized images. These analyses prove the usefulness of matrices when working with large sets of finite data, and they merely serve as an introduction to the powerful information processing that these tools provide.

11 Code

The entire codebase for the project follows, and is available for download  here.¹

11.1 Python

The Python code to generate the images is included below.

```

1  #!/usr/bin/env python
2  '''
3  APPM 2360 Differential Equations Project Two
4      |-Will Farmer
5      |-Jeffrey Milhorn
6      |-Patrick Harrington
7
8  This code takes the two given images and performs several
9  mathematical operations on them using matrix methods.
10 '''
11
12 import sys                                # Import system library
13 import scipy.misc                         # Import image processing libraries
14 import numpy                             # Import matrix libraries
15 import matplotlib.pyplot as plt          # Import plotting libraries
16 import pp                                 # Library for Parallel Processing
17
18 jobServer = pp.Server() # Create a new jobserver
19 jobs      = []          # List of jobs to complete
20
21 def main():
22     # Open images for manipulation
23     print('Opening Images')
24     image1 = scipy.misc.imread('../img/photo1.jpg')
25     image2 = scipy.misc.imread('../img/photo2.jpg')
26     image3 = scipy.misc.imread('../img/sadfox.jpg')
27
28     # Run manipulations on both images
29     print('Generating Manipulations')
30     manipulate(image1, '1')
31     manipulate(image2, '2')
32     manipulate(image3, '3')
33
34     # Visualize Determinants of DST Matrix
35     print('Generating Determinant Graph')
36     visualize_s()
37

```

¹ If you are unable to download these attached files, please go to [this link](#)

```

38     # Compress images using DST
39     print('Compressing Images')
40     jobs.append(
41         jobServer.submit(compression,
42                         (image1, '1', 0.5),
43                         (create_grayscale, dst, create_S),
44                         ('numpy', 'scipy.misc'))
45         ) # Add a new job to compress our first image
46     jobs.append(
47         jobServer.submit(compression,
48                         (image2, '2', 0.5),
49                         (create_grayscale, dst, create_S),
50                         ('numpy', 'scipy.misc'))
51         ) # Add a new job to compress our second image
52     jobs.append(
53         jobServer.submit(compression,
54                         (image3, '3', 0.5),
55                         (create_grayscale, dst, create_S),
56                         ('numpy', 'scipy.misc'))
57         ) # Add a new job to compress our third image
58
59     # Analyze Compression Effectiveness
60     print('Generating Compression Effectiveness')
61     comp_effect(image1, image2, image3)
62
63     # Create Picture Grid
64     print('Generating Picture Grid')
65     mass_pics(image1, '1')
66     mass_pics(image2, '2')
67     mass_pics(image3, '3')
68
69     for job in jobs:
70         job() # Evaluate all current jobs
71     jobServer.get_stats()
72
73 def manipulate(image, name):
74     '''
75     Manipulate images as directed
76     1) Create grayscale image
77     2) Produce horizontal shifts
78     3) Produce Vertical/Horizontal Shifts
79     4) Flip image vertically
80     '''
81     # Create grayscale
82     g = create_grayscale(image.copy())
83     scipy.misc.imsave('../img/gray%s.png' %name, g)
84
85     # Shift Horizontally
86     hs = shift_hort(g)
87     scipy.misc.imsave('../img/hsg%s.png' %name, hs)
88
89     # Shift Hort/Vert
90     hs = shift_hort(g)
91     vhs = shift_vert(hs.copy())

```

```

92     scipy.misc.imsave('../img/vhsg%s.png' %name, vhs)
93
94     # Flip
95     flipped = flip(g)
96     scipy.misc.imsave('../img/flip%s.png' %name, flipped)
97
98 def flip(image):
99     '''
100     flips an image
101     Essentially just multiplies it by a flipped id matrix
102     '''
103     il = numpy.identity(len(image)).tolist() # Creates a matching identity
104     for row in il: # Reverses the identity matrix
105         row.reverse()
106     i = numpy.array(il) # Turns it into a formal array
107     return numpy.dot(i, image) # Dots them together
108
109 def shift_hort(image):
110     '''
111     Shift an image horizontally
112     1) Create rolled identity matrix:
113         | 0 0 1 |
114         | 1 0 0 |
115         | 0 1 0 |
116     2) Dot with image
117     '''
118     i = numpy.roll(numpy.identity(len(image[0])),
119                    240, axis=0) # Create rolled idm
120     shifted = numpy.dot(image, i) # dot with image
121     return shifted
122
123 def shift_vert(image):
124     '''
125     Shift an image horizontally
126     1) Create rolled identity matrix:
127         | 0 0 1 |
128         | 1 0 0 |
129         | 0 1 0 |
130     2) Dot with image
131     '''
132     i = numpy.roll(numpy.identity(len(image)),
133                    100, axis=0) # create rolled idm
134     shifted = numpy.dot(i, image) # dot with image
135     return shifted
136
137 def create_grayscale(image):
138     '''
139     Creates grayscale image from given matrix
140     1) Create ratio matrix
141     2) Dot with image
142     '''
143     ratio = numpy.array([30., 59., 11.])
144     return numpy.dot(image.astype(numpy.float), ratio)
145

```

```

146 def shift_hort_color(image):
147     '''
148     Shift a color image horizontally
149     1) Create identity matrix that looks as such:
150         | 0 0 1 |
151         | 1 0 0 |
152         | 0 1 0 |
153     2) Dot it with image matrix
154     3) Return Transpose
155     '''
156     # Create an identity matrix and roll the rows
157     i = numpy.roll(
158         numpy.identity(
159             len(image[0]))
160         , 240, axis=0)
161     shifted = numpy.dot(i, image) # Dot with image
162     return numpy.transpose(shifted) # Return transpose
163
164 def compression(image, name, p):
165     '''
166     Compress the image using DST
167     '''
168     g = create_grayscale(image.copy()) # Create grayscale image matrix copy
169     t = dst(g) # Acquire DST matrix of image
170     (row_size, column_size) = numpy.shape(t) # Size of t
171     for row in range(row_size):
172         for col in range(column_size):
173             if (row + col + 2) > (2 * p * column_size):
174                 t[row][col] = 0 # if the data is above a set line, delete it
175     scipy.misc.imsave('../img/comp%s.png' %name, dst(t))
176
177 def dst(image):
178     '''
179     If given a grayscale image array, use the DST formula
180     and return the result
181     Uses this method:
182         image = X
183         DST = S
184         Y = S.(X.S)
185     '''
186     rows = numpy.dot(image, create_S(len(image[0])))
187     columns = numpy.dot(create_S(len(image)), rows)
188     return columns
189
190 def create_S(n):
191     '''
192     Discrete Sine Transform
193     1) Initialize variables
194     2) For each row and column, create an entry
195     '''
196     new_array = [] # What we will be filling
197     size = n
198     for row in range(size):
199         new_row = [] # New row for every row

```

```

200     for col in range(size):
201         S = ((numpy.sqrt(2.0 / size)) * # our equation
202             (numpy.sin((numpy.pi * ((row + 1) - (1.0/2.0)) *
203                 ((col + 1) - (1.0/2.0)))/(size))))
204         new_row.append(S) # Append entry to row list
205     new_array.append(new_row) # append row to array
206     return_array = numpy.array(new_array)
207     return return_array
208
209 def mass_pics(image, name):
210     '''
211     Create a lot of compressed Pictures
212     '''
213     answer = raw_input('Create .gif Images? (y/n) ')
214     if answer == 'n':
215         return None # It takes a while, so it's optional
216     domain = numpy.arange(0, 1.01, 0.01) # Range of p vals
217     for p in domain:
218         jobs.append(
219             jobServer.submit(compression,
220                             (image, 'array_%s_%f' %(name, p), p),
221                             (create_grayscale, dst, create_S),
222                             ('numpy', 'scipy.misc'))
223             ) # For each value of p, add a new compression job
224
225 def visualize_s():
226     '''
227     DST
228     Visualize the discrete sine transform equation implemented below.
229     Uses matplotlib to create graph
230     '''
231     nrange = numpy.arange(1, 33, 1) # Create values range [1,32] stepsize 1
232     det_plot = plt.figure() # New matplotlib class instance for a figure
233     det_axes = det_plot.add_axes([0.1, 0.1, 0.8, 0.8]) # Add axes to figure
234     yrange = [] # Create an empty y range (we'll be adding to this)
235     for number in nrange:
236         array = create_S(number) # Get a new array with size n
237         yrange.append(numpy.linalg.det(array)) # append determinant to yrange
238     det_axes.plot(nrange, yrange, label='Set of determinants') # Create line
239     det_axes.plot(nrange, yrange, 'ro') # Add points
240     det_axes.plot(nrange, nrange*0, 'k:') # Also create line at y=0
241     det_axes.legend(loc=4) # Place legend
242     plt.xlabel('Size of Discrete Sine Transform Matrix') # Label X
243     plt.ylabel('Determinant of Matrix') # Label Y
244     plt.title('Size of Matrix vs. its Determinant') # Title
245     plt.savefig('../img/dst_dets.png') # Save as a png
246
247 def comp_effect(image1, image2, image3):
248     '''
249     Analyzes compression effectiveness
250     If the image already exists, it will not run this
251     '''
252     try:
253         open('../img/bitcount.png', 'r')

```

```

254     open('../img/bitrat.png', 'r')
255     print(' |-> Graphs already created, skipping.\
256           (Delete existing graphs to recreate)')
257     # If it already exists, don't create it. (It takes a while)
258 except IOError:
259     g1 = create_grayscale(image1.copy()) # Create grayscale from copy of 1
260     g2 = create_grayscale(image2.copy()) # Create grayscale from copy of 2
261     g3 = create_grayscale(image3.copy()) # Create grayscale from copy of 2
262
263     domain1 = numpy.arange(0.0, 1.01, 0.01) # Range of p values
264     domain2 = numpy.arange(0.0, 1.01, 0.01) # Range of p values
265     domain3 = numpy.arange(0.0, 1.01, 0.01) # Range of p values
266
267     # Parallelize System and generate range
268     local_jobs = []
269     local_jobs.append(jobServer.submit(get_yrange,
270                                       (domain1, g1),
271                                       (dst, clear_vals, create_S),
272                                       ('numpy', 'scipy.misc')))
273     local_jobs.append(jobServer.submit(get_yrange,
274                                       (domain2, g2),
275                                       (dst, clear_vals, create_S),
276                                       ('numpy', 'scipy.misc')))
277     local_jobs.append(jobServer.submit(get_yrange,
278                                       (domain3, g3),
279                                       (dst, clear_vals, create_S),
280                                       ('numpy', 'scipy.misc')))
281
282     results = []
283     for job in local_jobs:
284         results.append(job())
285     count_y1 = results[0][0] # Assign variables
286     rat_y1   = results[0][1]
287     count_y2 = results[1][0]
288     rat_y2   = results[1][1]
289     count_y3 = results[2][0]
290     rat_y3   = results[2][1]
291
292     count_plot = plt.figure() # New class instance for a figure
293     count_axes = count_plot.add_axes([0.1, 0.1, 0.8, 0.8]) # Add axes
294     count_axes.plot(domain1, count_y1, label='Image 1')
295     count_axes.plot(domain2, count_y2, label='Image 2')
296     count_axes.plot(domain3, count_y3, label='Image 3')
297     count_axes.legend(loc=4)
298     plt.xlabel("Value of p")
299     plt.ylabel("Number of Non-Zero Bytes")
300     plt.title("Compression Effectiveness")
301     plt.savefig("../img/bitcount.png")
302
303     ratio_plot = plt.figure() # New class instance for a figure
304     ratio_axes = ratio_plot.add_axes([0.1, 0.1, 0.8, 0.8]) # Add axes
305     ratio_axes.plot(domain1, rat_y1, label='Image 1')
306     ratio_axes.plot(domain2, rat_y2, label='Image 2')
307     ratio_axes.plot(domain3, rat_y3, label='Image 3')
308     ratio_axes.legend(loc=4)

```



```
308     plt.xlabel("Value of p")
309     plt.ylabel("Ratio of Non-Zero Bytes to Total Bytes")
310     plt.title("Compression Effectiveness")
311     plt.savefig("../img/bitrat.png")
312
313 def get_yrange(domain, g):
314     bit_count = [] # Range for image
315     bit_ratio = []
316     for p in domain:
317         t = dst(g.copy()) # Transform 1
318         initial_count = float(numpy.count_nonzero(t))
319         clear_vals(t, p) # Strip of high-freq data
320         final_count = float(numpy.count_nonzero(t))
321         bit_count.append(final_count) # Append number of non-zero entries
322         bit_ratio.append(final_count / initial_count)
323     return bit_count, bit_ratio
324
325 def clear_vals(transform, p):
326     '''
327     Takes image and deletes high frequency
328     '''
329     (row_size, column_size) = numpy.shape(transform) # Size of t
330     for row in range(row_size):
331         for col in range(column_size):
332             if (row + col + 2) > (2 * p * column_size):
333                 transform[row][col] = 0 # if the data is above line, delete it
334     return transform
335
336 if __name__ == '__main__':
337     sys.exit(main())
```

11.2 MATLAB Code

Some MATLAB Code was also made that features equivalent functionality

Grayscale

```

1 function gray_image=grayscale(image)
2 % This is a function to take an image in jpg form and put it into grayscale
3
4 % This reads in the image
5 image_matrix=imread(image);
6
7 % get the dimensions
8 [rows,columns,~]=size(image_matrix);
9
10 % preallocate
11 gray_image = zeros(rows,columns);
12 for a=1:rows;
13     for b=1:columns;
14         gray_image(a,b)=0.3*image_matrix(a,b,1)...
15             +0.59*image_matrix(a,b,2)...
16             +0.11*image_matrix(a,b,3);
17     end
18 end
19 imwrite(uint8(gray_image),'name.jpg')
20
21 end

```

Horizontal Shifting

```

1 function [hshifted_image] = hshift(image)
2
3 % c is the number of cols we want to shift by
4 c = 240;
5
6 % read in the image and make it a nice little matrix
7 image_matrix=double(imread(image));
8
9 % get the dimensions of the matrix
10 [rows, cols] = size(image_matrix);
11
12 % get the largest dimension for the identity matrix
13 n = max(rows, cols);
14
15 % Preallocate for the id matrix:
16 T = zeros(n,n);
17
18 % generate a generic identity matrix
19 id = eye(n);
20
21 %fill in the first c cols of T with the last c cols of id
22 T(:,1:c)=id(:,n-(c-1):n);
23 %fill in the rest of T with the first part of id

```

```

24 T(:,c+1:n) = id(:,1:n-c);
25
26 hshifted_image=uint8(image_matrix*T);
27
28 imwrite(hshifted_image,'hshifted.jpg');

```

Vertical Shifting

```

1 function [vshifted_image] = vshift(image)
2
3 % r is the number of rows we want to shift by
4 r = 100;
5
6 % read in the image and make it a nice little matrix
7 image_matrix=double(imread(image));
8
9 % get the dimensions of the matrix
10 [rows, cols] = size(image_matrix);
11
12 % get the largest dimension for the identity matrix
13 n = min(rows, cols);
14
15 % Preallocate for the id matrix:
16 T = zeros(n,n);
17
18 % generate a generic identity matrix
19 id = eye(n);
20
21 %fill in the first c cols of T with the last c cols of id
22 T(1:r,:) = id(n-(r-1):n,:);
23 %fill in the rest of T with the first part of id
24 T(r+1:n,:) = id(1:n-r,:);
25
26 vshifted_image=uint8(T*image_matrix);
27
28 imwrite(vshifted_image,'vshifted.jpg');

```

11.3 MATLAB Code

Some MATLAB Code was also made that features equivalent functionality

Grayscale

```

1 function gray_image=grayscale(image)
2 % This is a function to take an image in jpg form and put it into grayscale
3
4 % This reads in the image
5 image_matrix=imread(image);
6
7 % get the dimensions
8 [rows,columns,~]=size(image_matrix);
9

```

```

10 % preallocate
11 gray_image = zeros(rows,columns);
12 for a=1:rows;
13     for b=1:columns;
14         gray_image(a,b)=0.3*image_matrix(a,b,1)...
15             +0.59*image_matrix(a,b,2)...
16             +0.11*image_matrix(a,b,3);
17     end
18 end
19 imwrite(uint8(gray_image),'name.jpg')
20
21 end

```

Horizontal Shifting

```

1 function [hshifted_image] = hshift(image)
2
3 % c is the number of cols we want to shift by
4 c = 240;
5
6 % read in the image and make it a nice little matrix
7 image_matrix=double(imread(image));
8
9 % get the dimensions of the matrix
10 [rows, cols] = size(image_matrix);
11
12 % get the largest dimension for the identity matrix
13 n = max(rows, cols);
14
15 % Preallocate for the id matrix:
16 T = zeros(n,n);
17
18 % generate a generic identity matrix
19 id = eye(n);
20
21 %fill in the first c cols of T with the last c cols of id
22 T(:,1:c)=id(:,n-(c-1):n);
23 %fill in the rest of T with the first part of id
24 T(:,c+1:n) = id(:,1:n-c);
25
26 hshifted_image=uint8(image_matrix*T);
27
28 imwrite(hshifted_image,'hshifted.jpg');

```

Vertical Shifting

```

1 function [vshifted_image] = vshift(image)
2
3 % r is the number of rows we want to shift by
4 r = 100;
5
6 % read in the image and make it a nice little matrix

```

```
7 image_matrix=double(imread(image));
8
9 % get the dimensions of the matrix
10 [rows, cols] = size(image_matrix);
11
12 % get the largest dimension for the identity matrix
13 n = min(rows, cols);
14
15 % Preallocate for the id matrix:
16 T = zeros(n,n);
17
18 % generate a generic identity matrix
19 id = eye(n);
20
21 %fill in the first c cols of T with the last c cols of id
22 T(1:r,:) = id(n-(r-1):n,:);
23 %fill in the rest of T with the first part of id
24 T(r+1:n,:) = id(1:n-r,:);
25
26 vshifted_image=uint8(T*image_matrix);
27
28 imwrite(vshifted_image,'vshifted.jpg');
```